

Identifying (deep) fake news communities on multiple social networks

Thiago Amado Costa¹, Humberto Torres Marques Neto¹

¹ICEI – Pontifícia Universidade Católica de Minas Gerais (PUC-MG)
Belo Horizonte, Minas Gerais - Brasil

thiago.amado@sga.pucminas.br, humberto@sga.pucminas.br

Abstract. *The misuse of Deepfakes and the rise of Fake News have become a significant challenge in the era of social media, posing serious threats to the credibility of information shared online. Additionally, individuals and groups can become targets of these malicious actions, creating communities that believe and share news of the same malicious sources. This problem is not exclusive to a single social media platform, which highlights the urgent need for robust fact-checking mechanisms, as well as solutions to identify these specific communities, that will help create a safer online environment.*

1. Introduction

Fake News and Deepfakes are one of the biggest challenges in today's social media, threatening the credibility of information shared online. Fake News is defined as incorrect or misleading information fabricated to mimic the structure of the news media, but without the process of ensuring accuracy and credibility ([Lazer et al. 2018]), and Deepfake is a form of multimedia manipulation that leverages advanced machine learning and artificial intelligence techniques to manipulate or generate visual and audio content, with a high capacity to deceive viewers ([Kietzmann et al. 2020]).

This phenomenon poses serious threats to today's social media, as it spreads faster and more broadly than credible information ([Vosoughi et al. 2018]), and Deepfakes are becoming more accessible and believable ([Kietzmann et al. 2020]). This reinforces the importance of further research on this topic. Hence, the primary objectives of this study are to analyze various methodologies for detecting Fake News and Deepfakes, as well as to examine and identify the communities associated with them.

2. Related Work

Identifying and mitigating the spread of fake news has become a pressing concern in recent years, as the proliferation of social media platforms has enabled the rapid dissemination of misinformation. In this section, we review existing research that focuses on detecting fake news, as well as communities across multiple social networks.

2.1. Fake News

The spread of fake news has become a major concern, impacting public opinion and civil debate. To combat this, researchers from many fields are developing methods to identify false information. This section explores key studies and approaches used in detecting fake news.

[Ruchansky et al. 2017] combines text, response, and source characteristics of news articles to approach the fake news detection. Unlike existing solutions that focus on individual aspects, the CSI model integrates these three key features using deep neural networks to automatically select important information and capture temporal dependencies in user engagement. By providing insights on both articles and users, CSI offers a more comprehensive understanding of fake news propagation. Experimental results on real-world datasets demonstrate that the CSI model achieves higher classification accuracy compared to previous methods, while also requiring fewer parameters and training efforts. This highlights the effectiveness of the CSI model in accurately detecting fake news and addressing the challenges associated with automated fake news detection.

[Kaliyar et al. 2021] introduces a new model called DeepFakeE for improving fake news detection by combining news content and social context. The approach combines news content and social context-based features in a tensor, resulting in a higher-order decision boundary and achieving state-of-the-art results for fake news detection. The article shows the significance of examining the social context for improving fake news detection. It emphasizes the role of user social engagements and echo chambers in the dissemination of fake news. Echo chambers, which are groups of like-minded individuals, play a vital role in the spread of fake news as they tend to share stories that align with their beliefs. The article suggests that utilizing echo chambers to obtain user-related and community-related information from news articles can enhance fake news detection.

[Chandra et al. 2020] discusses the use of different baselines for fake news detection, including text-based, majority sharing, and social baselines. It compares the article sharing behavior of users in two datasets, GossipCop and HealthStory, highlighting differences in user characteristics and community graph structures. The study shows that GNNs are more effective in detecting fake articles in GossipCop due to richer community-level features, while struggling in HealthStory due to limited user representations and mixed article sharing patterns. The proposed graph-based approach outperforms text-based models, emphasizing the importance of community-based modeling in fake news detection.

[Shu et al. 2018] explores the role of social context in detecting fake news, proposing a TriFN framework that considers the tri-relationship among publishers, news pieces, and users for classification. It highlights that social context features are more effective than news content features in predicting fake news, with user-based, post-based, and network-based features playing crucial roles. The study demonstrates that the TriFN approach significantly outperforms other baseline methods for fake news detection, showcasing the importance of considering social context in combating fake news.

2.2. Social Media Communities

Community detection has garnered significant attention across various disciplines due to its applications in understanding the structure and dynamics of complex systems. Over the years, researchers have proposed numerous algorithms and methods to identify cohesive groups within networks. In this section, we provide an overview of key works in the field.

[Raghavan et al. 2007] proposes a new algorithm for efficiently detecting communities in large networks. The algorithm leverages a label propagation approach where nodes iteratively adopt the most common label among their neighbors. This simple strategy facilitates the formation of communities without requiring prior knowledge about

their structure or size. The effectiveness of the algorithm is demonstrated on well-established benchmarks, achieving high accuracy. Moreover, the algorithm boasts a near-linear time complexity, making it significantly faster than existing methods. Furthermore, it bypasses the need for complex function optimization or parameter tuning, offering a more straightforward and computationally efficient approach to community detection in large-scale networks.

[Zhang et al. 2016] proposes a parallel label propagation algorithm to identify communities within social networks. It tackles shortcomings of the traditional label propagation method. The new algorithm incorporates grey relational analysis during label updates, leading to less randomness in label selection. Additionally, it optimizes the parallel computation steps. Experiments on various social networks, both simulated and real-world, demonstrate that this approach is scalable and delivers high accuracy in identifying communities

[Chunaev 2020] dives into the challenge of uncovering communities within social networks, highlighting the importance of considering both the network's structure (connections between users) and the attributes of individual users (interests, demographics). A key point addressed is the current lack of knowledge on how these two aspects influence the accuracy of community detection. To shed light on this, the article provides a classification system for different community detection methods based on how they integrate network structure and node attributes. This classification goes beyond just naming categories, offering technical details and insights into how these methods perform. The article explores various approaches like those based on metaheuristics, probabilistic models, and reaching consensus among different detection strategies.

[Zhao et al. 2021] introduces a community detection algorithm (CDEP) designed for massive social networks. The algorithm leverages graph compression techniques to efficiently identify communities within these complex structures. Notably, CDEP outshines several existing cutting-edge algorithms in terms of both effectiveness and efficiency. When tested on various social networks, CDEP consistently revealed community structures that closely matched real-world groupings. Moreover, it achieved optimal modularity scores, a key metric for community detection quality. Furthermore, CDEP stands out as the only method capable of tackling large-scale networks within a practical time frame, demonstrating its exceptional scalability.

[Cinelli et al. 2021] investigates how social media platforms like Facebook, Twitter, and Reddit contribute to the echo chamber effect. The researchers define ways to measure echo chambers by examining how users with similar viewpoints connect and share information (homophily in interaction networks) and how information slants towards reinforcing existing beliefs (bias in information diffusion). They also establish methods for pinpointing user attitudes on specific topics within these platforms. By comparing information flow across different social media, the study reveals how content tends to spread more readily among users who already share similar perspectives.

In [Ruiz and Nilsson 2023], social media is portrayed as a breeding ground for disinformation that thrives on identity politics. Actors behind the spread of disinformation weaponize rhetoric to manipulate grievances and group affiliation. This manipulation isn't passive; consumers actively engage in arguments, promoting narratives that challenge es-

tablished facts. The article highlights the echo chamber effect as a culprit in solidifying these viewpoints. Through constant debate within closed communities, arguments become ingrained in personal identities. The example of Flat Earthers who leverage biblical interpretations to justify their beliefs exemplifies this phenomenon, where they position themselves as defenders of a purer form of Christianity.

References

- Chandra, S., Mishra, P., Yannakoudakis, H., Nimishakavi, M., Saeidi, M., and Shutova, E. (2020). Graph-based modeling of online communities for fake news detection.
- Chunaev, P. (2020). Community detection in node-attributed social networks: A survey. *Computer Science Review*, 37:100286.
- Cinelli, M., Morales, G., Galeazzi, A., Quattrociocchi, W., and Starnini, M. (2021). The echo chamber effect on social media. *Proceedings of the National Academy of Sciences*, 118:e2023301118.
- Kaliyar, R., Goswami, A., and Narang, P. (2021). Deepfake: improving fake news detection using tensor decomposition-based deep neural network. *The Journal of Supercomputing*, 77.
- Kietzmann, J., Lee, L. W., McCarthy, I. P., and Kietzmann, T. C. (2020). Deepfakes: Trick or treat? *Business Horizons*, 63(2):135–146. ARTIFICIAL INTELLIGENCE AND MACHINE LEARNING.
- Lazer, D. M., Baum, M. A., Benkler, Y., Berinsky, A. J., Greenhill, K. M., Menczer, F., Metzger, M. J., Nyhan, B., Pennycook, G., Rothschild, D., et al. (2018). The science of fake news. *Science*, 359(6380):1094–1096.
- Raghavan, U. N., Albert, R., and Kumara, S. (2007). Near linear time algorithm to detect community structures in large-scale networks. *Physical Review E*, 76(3).
- Ruchansky, N., Seo, S., and Liu, Y. (2017). Csi: A hybrid deep model for fake news detection. In *Proceedings of the 2017 ACM on Conference on Information and Knowledge Management, CIKM '17*. ACM.
- Ruiz, C. D. and Nilsson, T. (2023). Disinformation and echo chambers: How disinformation circulates on social media through identity-driven controversies. *Journal of Public Policy & Marketing*, 42(1):18–35.
- Shu, K., Wang, S., and Liu, H. (2018). Beyond news contents: The role of social context for fake news detection.
- Vosoughi, S., Roy, D., and Aral, S. (2018). The spread of true and false news online. *Science*, 359(6380):1146–1151.
- Zhang, Q., Qiu, Q., Guo, W., Guo, K., and Xiong, N. (2016). A social community detection algorithm based on parallel grey label propagation. *Computer Networks*, 107:133–143. Machine learning, data mining and Big Data frameworks for network monitoring and troubleshooting.
- Zhao, X., Liang, J., and Wang, J. (2021). A community detection algorithm based on graph compression for large-scale social networks. *Information Sciences*, 551:358–372.